

The Explanation Is the Forward Pass: Interpretability by Construction with Certified Legible Bottlenecks

Anwar Debes*

University of Agder, Norway
anward@uia.no

July 2026

Abstract

The dominant paradigm in interpretability is *archaeological*: train an opaque model, then dig for structure after the fact with probes, sparse autoencoders (SAEs), and attribution graphs. Each of these methods measures faithfulness *post hoc*; none *guarantees* it, and by 2026 the limits are well documented: SAEs do not recover canonical features and often fail to beat linear probes, attribution graphs leave unbounded computation in “error nodes”, and chain-of-thought is an unfaithful, optimisation-fragile window. We argue that interpretability should instead be an *architectural invariant the model is required to satisfy*, with a guarantee emitted at inference time. We introduce the **Legible Bottleneck**, a single composable primitive that reads a layer’s computation through a sparse, typed, persistently named concept code and reconstructs it; the unexplained remainder, the *leak*, is kept and measured rather than discarded. Stacking Legible Bottlenecks yields **CANDOR** (Concept-ANchored Disentangled Output Routing), a network co-trained so that the named concepts reconstruct the computation, are *causally sufficient* (interchange interventions as a training signal, not a post-hoc test), and keep anchored identities across training. On every forward pass, CANDOR emits a **Faithfulness Certificate** $\delta(x)$: the total-variation distance between the deployed model and its leak-ablated “legible” replacement. Because δ is a measurement rather than a learned monitor, it is sound by construction and re-checkable by any third party. On synthetic tasks with known ground-truth concepts and mechanism, CANDOR recovers the planted concepts (mean correlation 0.919), certifies its decisions with mean $\delta = 0.0021$ at no measurable interpretability tax, and passes held-out causal tests (81.7% directional accuracy); a backdoor that bypasses the concept channel is exposed by a certificate spike (detection AUC 0.901). The primitive composes with attention. On the pretrained GPT-2 the certificate machinery transfers, but neither a *frozen* retrofit nor single-layer fine-tuning favours CANDOR over a reconstruction-only SAE, and an $\approx 8\times$ reconstruction improvement under fine-tuning fails to tighten the certificate: reconstruction and faithfulness dissociate, and full by-construction training remains the central open question. We are explicit about scope: these are small-scale demonstrations, and completeness is unattainable under computation in superposition; the contribution is to make the bypass *measured, bounded, and certified*. All code and every measured number are released.

Keywords: mechanistic interpretability; sparse autoencoders; concept bottleneck models; causal abstraction; faithfulness; certified explanations; AI safety.

*Final-year student, Master’s programme in Artificial Intelligence, University of Agder, Norway. Correspondence: anward@uia.no. Code, the reference implementation, and every measured number are available at <https://github.com/AnwarDebes/candor>.

1 Introduction

For a decade interpretability has been practised as *archaeology*: train a model whose computation is opaque, and then, with linear probes, sparse autoencoders, activation patching, attribution graphs, and chain-of-thought inspection, excavate. The excavations are remarkable: induction heads [Olsson et al., 2022], a circuit for indirect object identification [Wang et al., 2023], millions of monosemantic features [Templeton et al., 2024], and per-prompt attribution graphs that traced planning and multi-step reasoning in a production model [Lindsey et al., 2025]. But archaeology has a structural problem that scaling does not fix: it produces *descriptions* of a model, validated *after the fact*, never a *guarantee* that the description is what the model computes.

By 2026 the limits are no longer speculative. Sparse autoencoders (SAEs), the field’s workhorse, were shown not to find canonical units of analysis [Leask et al., 2025], to suffer feature absorption and hedging [Chanin et al., 2024, 2025], and, decisively, to *not* reliably beat linear probes or even random baselines on the supervised tasks one would want them for [Kantamneni et al., 2025]. Attribution graphs, the state of the art in tracing computation, route an admitted and unbounded fraction of the work through “error nodes” that the explanation simply does not cover [Ameisen et al., 2025]. Chain-of-thought, the most accessible window of all, is unfaithful [Turpin et al., 2023, Chen et al., 2025] and, worse, becomes *actively misleading* under optimisation pressure [Baker et al., 2025, Korbak et al., 2025]. The common root is that every one of these methods measures faithfulness post hoc; none makes the model *be* faithful.

Thesis. An advanced agent whose decisions must be trusted cannot be audited by digging. We propose the opposite stance: make legibility a property the architecture is *required* to have, and emit the audit *at inference time*. Concretely, force the output-relevant computation through a sparse, typed, persistently named concept channel, co-train that channel to be *causally* faithful, and measure, on every forward pass, how much computation escaped it. The escaped fraction is not hidden; it is the certificate. This reframes the central question from “can one find structure in the model?” to “how much of the model’s decision is *not* accounted for by named concepts, and can the model prove a bound on that number right now?”

The primitive. The central object is the **Legible Bottleneck** (§3): at a checkpoint with activation h , encode a sparse non-negative code $c = \text{TopK}(\text{ReLU}(W_{\text{enc}}(h - b_{\text{dec}})))$, decode $\hat{h} = W_{\text{dec}}c + b_{\text{dec}}$, and route the network with the *leak* $r = h - \hat{h}$ kept and measured rather than thrown away. It is one operation, it composes, and, like self-attention [Vaswani et al., 2017] for sequence modelling, the architecture is just the primitive, stacked. We call the stack **CANDOR** (Concept-ANchored Disentangled Output Routing, §4).

The guarantee. The deployed model M has a twin, the *legible replacement* M_{leg} , obtained by ablating every leak; M_{leg} ’s output is a pure function of the named concept codes. The **Faithfulness Certificate** is $\delta(x) = \text{TV}(M(x), M_{\text{leg}}(x))$, the total-variation distance between the deployed model and its named-concept explanation (§6). Crucially δ is *measured*, not predicted by a learned head, so it cannot be Goodharted into silence the way a learned monitor can [Baker et al., 2025]: it is a sound, one-sided, per-input, re-checkable upper bound on the influence of all non-legible (“dark”) computation.

Contributions.

1. **The Legible Bottleneck** (§3) and the **CANDOR** architecture (§4): a by-construction realisation of “the explanation is the forward pass” in which output-relevant computation is routed through a sparse, typed, anchored concept channel and the unexplained remainder is a first-class, measured quantity.

2. **A training objective (§5)** that conjoins completeness, *causal* faithfulness (interchange and causal-scrubbing interventions as a training signal), and concept anchoring; we show empirically that the conjunction is necessary: removing any term makes the certificate vacuous or the channel leak.
3. **The Faithfulness Certificate (§6)**: a per-forward-pass, sound-by-construction, re-checkable bound on dark computation, with the precise predicate it does and does not guarantee stated honestly (Propositions 1–3).
4. **A ground-truth evaluation, and an honest real-LLM probe (§7)**: on synthetic tasks with known concepts and known mechanism, where soundness can actually be checked, CANDOR recovers the concepts, certifies tightly at near-zero tax, passes a directional causal test, detects a channel-bypassing backdoor via a certificate spike, and composes with attention on a sequence task. On the pretrained GPT-2 we show the certificate machinery transfers to a real LLM *and*, honestly, that neither a frozen retrofit nor single-layer fine-tuning favours CANDOR over an SAE, while an $\approx 8\times$ reconstruction improvement fails to tighten the certificate, dissociating reconstruction from faithfulness and delimiting the contribution to the by-construction regime. Every number is measured and released.

We do not claim to have solved interpretable AI. We claim a concrete, checkable mechanism that converts interpretability from an after-the-fact description into an architectural invariant with a runtime guarantee, together with an honest account of exactly how far that guarantee currently reaches.

2 Background: six stages toward legibility

The modern interpretability literature reads as a sequence of attempts to make neural computation legible, each born from the failure of the last; CANDOR sits at the end of the arc and is best understood as a conjunction of its final three stages.

Stage 0: the substrate (superposition). Networks pack many more sparse, near-linear features than they have neurons [Elhage et al., 2022, 2021, Olah et al., 2020], which both motivates a sparse legible basis and bounds it: if *computation* (not just storage) occurs in superposition [Hänni et al., 2024], no fixed-width legible channel can be complete. We concede this up front.

Stage 1: recover features post hoc (SAEs). Dictionary learning disentangles superposition into nameable features [Bricken et al., 2023, Cunningham et al., 2023, Templeton et al., 2024, Gao et al., 2024], with architectural refinements [Rajamanoharan et al., 2024, Bussmann et al., 2024, 2025] and open suites [Lieberum et al., 2024]. The 2025 crisis [Leask et al., 2025, Kantamneni et al., 2025, Chanin et al., 2024] diagnoses the flaw CANDOR targets: *no SAE objective is causally faithful by construction*, and reconstruction error is large and behaviourally load-bearing.

Stage 2: trace computation (transcoders, attribution graphs). Transcoders approximate a layer’s input-output map [Dunefsky et al., 2024]; cross-layer transcoders plus replacement models yield per-prompt attribution graphs [Ameisen et al., 2025, Lindsey et al., 2025, Marks et al., 2025]. This is the closest prior art to “the explanation is the forward pass.” Its limits (unbounded error nodes, per-prompt cost, post-hoc construction) are exactly what a runtime certificate must fix.

Stage 3: make faithfulness causal. The field converges on intervention as the only credible test: activation/path patching [Goldowsky-Dill et al., 2023, Zhang and Nanda, 2024], automated

discovery [Conmy et al., 2023, Syed et al., 2023], causal scrubbing [Chan et al., 2022], and the unifying theory of causal abstraction [Geiger et al., 2021, 2025]. Decisively, faithfulness can be a *training* objective: interchange intervention training (IIT) [Geiger et al., 2022] and distributed alignment search [Geiger et al., 2024, Wu et al., 2023]. This is CANDOR’s core borrowed engine.

Stage 4: build legibility in. Force computation through a privileged basis: concept bottleneck models [Koh et al., 2020, Espinosa Zarlenga et al., 2022], self-explaining networks [Alvarez-Melis and Jaakkola, 2018], prototypes [Chen et al., 2019], B-cos alignment [Böhle et al., 2022], codebooks [van den Oord et al., 2017], the information bottleneck [Tishby and Zaslavsky, 2015], and the call to use inherently interpretable models for high-stakes decisions [Rudin, 2019]. Their recurring failure is *concept leakage*: the bottleneck silently re-becomes a black box [Mahinpei et al., 2021, Havasi et al., 2022], and identifiability is unattainable without inductive bias [Locatello et al., 2019].

Stage 5: certify. Neural-network verification is mature for *robustness* [Katz et al., 2017, Zhang et al., 2018, Wang et al., 2021, Xu et al., 2020] but robustness is not interpretability; formal-XAI gives per-input logical guarantees [Ignatiev et al., 2019]; interpretable-by-construction-and-verified models exist only at small scale [Przybysz et al., 2023, Granmo, 2018, Lin et al., 2020]. No method emits a *sound, cheap, per-forward-pass* bound on bypassed computation.

CANDOR is the conjunction of Stages 3, 4 and 5 performed *jointly* and emitted *at runtime*. §8 states precisely what is borrowed from each and what is new.

3 The Legible Bottleneck

Let $h \in \mathbb{R}^d$ be the activation at a chosen checkpoint (e.g. an MLP hidden state). A Legible Bottleneck is parameterised by an overcomplete dictionary of $m \gg d$ atoms.

Definition 1 (Legible Bottleneck). *With encoder $W_{\text{enc}} \in \mathbb{R}^{m \times d}$, decoder $W_{\text{dec}} \in \mathbb{R}^{d \times m}$ (unit-norm columns), and biases $b_{\text{enc}}, b_{\text{dec}}$, define*

$$c = \text{TopK}_k(\text{ReLU}(W_{\text{enc}}(h - b_{\text{dec}}) + b_{\text{enc}})) \in \mathbb{R}_{\geq 0}^m, \quad \hat{h} = W_{\text{dec}} c + b_{\text{dec}}, \quad r = h - \hat{h}.$$

c is the concept code (at most k active named concepts), $\hat{h}$ the legible reconstruction, and r the leak.

The bottleneck does not replace h ; it *routes* the downstream network in one of three modes, and this routing, not the dictionary, is what separates the construction from a post-hoc sparse code:

- **full** (deployed): the network receives $h = \hat{h} + r$ unchanged. The bottleneck is an auxiliary read plus a regulariser, so capability is *never* destroyed by construction; the interpretability tax can only enter through training pressure, not through a lossy forward pass.
- **legible** (replacement): the network receives $\hat{h}$; the leak is ablated. The resulting model M_{leg} has an output that is a pure function of the named concept codes.
- **leak-swap** (causal probe): the network receives $\hat{h} + r_{\pi}$, the leak replaced by another input’s leak under a permutation π . If the named concepts are causally sufficient, the output is unchanged; the training objective drives this.

The TopK non-linearity (cf. TopK SAEs [Gao et al., 2024]) gives an exact $L_0 = k$ budget (a hard, legible cap on how many concepts may participate in a single decision), while unit-norm decoder atoms supply geometric identifiability. Each atom carries a persistent integer identity (a *name*) and an optional *type* tag; §5 keeps those identities stable across training.

4 The CANDOR architecture

A CANDOR model places Legible Bottlenecks at the checkpoints that carry output-relevant computation. In the reference implementation we study (i) an MLP whose concept layer sits on a frozen, identity-initialised embedding of the input (the site at which generative factors are linearly present, so the sparse code can disentangle them), followed by a task head, and (ii) a Transformer whose every MLP hidden state is a bottleneck site, demonstrating composition with attention. The deployed model M and its legible twin M_{leg} share all weights and differ only in routing mode, so M_{leg} is not a separate distilled model fitted after the fact (as a transcoder replacement model is [Ameisen et al., 2025]) but the *same network* read through its concept channel. The explanation of a decision is then literally the list of active named concepts and their causal contributions, available at inference with no search.

5 The training objective

CANDOR is trained with one loss whose terms are mutually enabling. Writing ℓ_{task} for cross-entropy, sites for the per-layer $(h, \hat{h}, r, c)$, and π for a random batch permutation,

$$\begin{aligned} \mathcal{L} = & \underbrace{\ell_{\text{task}}(M)}_{\text{accurate}} + \alpha \underbrace{\ell_{\text{task}}(M_{\text{leg}})}_{\text{legible model accurate}} + \beta \underbrace{\text{KL}(M \parallel M_{\text{leg}})}_{\text{faithfulness}} \\ & + \lambda \underbrace{\sum_{\ell} \frac{\|r^{(\ell)}\|^2}{\|h^{(\ell)}\|^2}}_{\text{leak energy}} + \mu \underbrace{\text{TV}(M, M_{\text{leak-swap}(\pi)})}_{\text{causal sufficiency}} + \gamma \underbrace{\sum_{\ell} \|W_{\text{dec}}^{(\ell)} - \bar{W}^{(\ell)}\|^2}_{\text{anchoring}}. \end{aligned}$$

(The implementation additionally applies a mild L_1 shrinkage to the concept codes; Appendix A lists every weight.) The *faithfulness* term trains the leak-ablated replacement to reproduce the deployed model’s behaviour, which drives the certificate δ toward zero. The *leak-energy* term pulls the activation onto the sparse-decodable manifold so the named code actually explains h . The *causal-sufficiency* term is the load-bearing, borrowed-from-IIT ingredient [Geiger et al., 2022, Chan et al., 2022]: it penalises the output’s sensitivity to swapping the (non-legible) leak across inputs, i.e. it trains the concepts to be a causal, not merely correlational, summary. The *anchoring* term penalises drift of each decoder atom from its persistent EMA identity $\bar{W}$, the inductive bias that identifiability theory proves necessary [Locatello et al., 2019] and that SAEs notably lack [Leask et al., 2025]. Objective ablations (§7) show each term is necessary: drop faithfulness or leak and δ becomes large (the channel leaks); drop causal and the leak regains causal power; drop anchoring and concept identities stop being stable across runs.

Listing 1: The Legible Bottleneck and the CANDOR step (reference implementation).

```
def legible_bottleneck(h, W_enc, b_enc, W_dec, b_dec, k):
    c = topk(relu(W_enc @ (h - b_dec) + b_enc), k) # sparse named concept code
    h_hat = W_dec @ c + b_dec # legible reconstruction
    return c, h_hat, h - h_hat # code, recon, leak

# one training step (3 forwards: deployed, legible, leak-swapped)
L_full = task(M(x, mode="full"), y)
L_leg = task(M(x, mode="legible"), y)
faith = KL(M(x, "full").detach(), M(x, "legible")) # replacement reproduces model
leak = mean(||r||^2 / ||h||^2 for sites) # shrink dark computation
causal = TV(M(x, "full").detach(), M(x, "full", leak_perm=randperm)) # leak is inert
loss = L_full + a*L_leg + b*faith + lam*leak + mu*causal + gam*anchor
```


6 The Faithfulness Certificate

Definition 2 (Faithfulness Certificate). *For input x , let $p_M(\cdot | x)$ and $p_{\text{leg}}(\cdot | x)$ be the output distributions of the deployed model and its legible replacement. The certificate is*

$$\delta(x) = \text{TV}(p_M(\cdot | x), p_{\text{leg}}(\cdot | x)) = \frac{1}{2} \sum_y |p_M(y | x) - p_{\text{leg}}(y | x)| \in [0, 1].$$

Proposition 1 (Soundness, by construction). *$p_{\text{leg}}(\cdot | x)$ depends on x only through the named concept codes $\{c^{(\ell)}(x)\}$. Hence for every event A over outputs, $|p_M(A | x) - p_{\text{leg}}(A | x)| \leq \delta(x)$: the named-concept explanation predicts the deployed model’s entire output distribution to within $\delta(x)$. Because $\delta(x)$ is computed by one extra (legible) forward pass rather than predicted by a learned head, it is exact and one-sided: it can never under-report behavioural divergence between the model and its explanation.*

Proof sketch. Total variation equals $\sup_A |p_M(A) - p_{\text{leg}}(A)|$, giving the bound for all A . M_{leg} ignores every leak by definition, so its output is a deterministic function of $\{c^{(\ell)}\}$. $\delta(x)$ is obtained by evaluating both forward passes and summing absolute differences; no parameter is fit to it, so it is a direct measurement. $\square$

Proposition 2 (Causal sufficiency $\Leftrightarrow$ leak-inertness). *Write the deployed output as $f(c, r)$ for concept code c and leak r . If for all inputs the output is invariant under replacing r by any other in-distribution leak r' (leak-swap invariance), then f does not depend on r on the data support, i.e. the output is determined by the named concepts alone. The causal term in $\mathcal{L}$ minimises exactly the leak-swap output deviation; the residual deviation and δ jointly quantify the failure of sufficiency.*

Proof sketch. For any two in-distribution leaks r, r' compatible with a code c , invariance gives $f(c, r) = f(c, r')$, so $f(c, \cdot)$ is constant on the support and the output is a function of c alone. The causal term is the expected output deviation under a random leak permutation, a Monte Carlo estimate of the violation of this invariance. $\square$

Proposition 3 (Checkability). *$\delta(x)$ and the explanation (active named concepts and, via concept ablation, their causal effects) are recomputable by any third party from the released weights with one extra forward pass per input, without access to training data or process. The certificate is therefore independently auditable.*

Proof sketch. Every quantity in the definitions above is a function of the released weights and the input alone: p_M and p_{leg} require one standard and one legible-mode forward pass, the active concepts are read off $c(x)$, and per-concept causal effects need one additional legible pass per ablated concept. $\square$

Remark 1 (No completeness claim). *CANDOR does not claim $\delta \equiv 0$ is attainable at scale. Computation in superposition [Hänni et al., 2024] implies some output-relevant work has no faithful sparse-linear description, so a non-trivial bypass can persist. The value of the certificate is soundness, not tightness: a large but honest $\delta(x)$ tells an auditor exactly when not to trust the named-concept explanation, which is precisely the signal post-hoc methods cannot provide. This also reframes the certificate as a constructive answer to eliciting latent knowledge [Christiano et al., 2021] and a quantitative mechanistic-anomaly detector [Roger et al., 2023]: dark computation is surfaced as a number, not hidden in an error node.*

7 Experiments

Following the discipline that interpretability claims are only meaningful where they can be checked, our primary evaluation is on synthetic data with known ground-truth concepts and a

Quantity	CANDOR	Opaque + post-hoc SAE
Task accuracy (deployed)	99.9%	99.9%
Task accuracy (legible / replacement)	99.9%	(n/a)
Interpretability tax	≈ 0 pp	(n/a)
Faithfulness certificate δ (mean)	0.0021	none emitted
model/explanation agreement	99.9%	(n/a)
Concept recovery (vs. ground truth)	0.919	0.966
Causal necessity gap	77.3 pp	(n/a)
Directional causal accuracy	81.7%	(n/a)
Non-legible remainder, causal effect	0.0054	0.0030
Per-input certificate	emitted	none

Table 1: Planted-concepts results. CANDOR attains concept recovery comparable to the standard train-then-excavate pipeline (a post-hoc SAE on the opaque model) at no measurable accuracy tax. The difference is *guarantees*: only CANDOR emits a per-input faithfulness certificate (δ), trains its concepts to be causally sufficient (the held-out necessity and directional tests), and, shown in the ablations, anchors their identity across runs. On this easily-reconstructed task the SAE’s uncaptured residual happens to be nearly inert too (0.0030 vs. CANDOR’s 0.0054), so incidental inertness is *not* the differentiator here; the certificate and the trained causal properties are. The regime where the certificate becomes decisive is computation that *bypasses* the channel (§7.2), which the post-hoc pipeline has no certificate to catch.

known mechanism. We then show composition with attention on a sequence task, and probe a real pretrained LLM (GPT-2). All numbers are measured by `experiments/run_all.py` and emitted to `paper/_numbers.tex`; nothing is hand-entered. Appendix A lists every architecture, hyperparameter, and seed.

7.1 Planted concepts: a checkable ground truth

We draw $K = 12$ independent sparse binary concepts, mix them linearly (in superposition) into a 48-dimensional observation, and set the label to the number of active concepts among a fixed relevant subset of size 5 (a 6-way task). Because the true concepts *and* the true mechanism are known, recovery, causal faithfulness, and certificate soundness are all directly checkable. We train a CANDOR-MLP (dictionary $m = 48$, $k = 8$, 2200 steps) and, for comparison, an opaque MLP of matched shape with a post-hoc TopK SAE fitted to its embedding, the standard “train-then-excavate” pipeline.

Table 1 summarises the headline result. CANDOR’s legible replacement solves the task at 99.9%, as accurate as the opaque baseline (99.9%), i.e. *no measurable interpretability tax*, while certifying every decision with a mean $\delta = 0.0021$ (95th percentile 0.0024) and agreeing with its own named-concept explanation 99.9% of the time. It recovers the 12 planted concepts (mean correlation 0.919, minimum 0.763); the post-hoc SAE recovers them slightly better (0.966). Recovery alone, however, does not differentiate the pipelines: both find the concepts. What CANDOR adds is everything that turns a *description* into a *guarantee*: a per-input certificate (the SAE pipeline emits none), causal sufficiency trained into the channel, anchored identity, and the ability to detect channel-bypassing computation (§7.2). On this clean task the SAE’s uncaptured residual is itself nearly inert (0.0030, comparable to CANDOR’s 0.0054), so incidental inertness is not what separates them here; the certificate and the trained causal properties are. Labelling activations after the fact can recover concepts, but it cannot certify that they are what the model used.

Causal faithfulness. On a *held-out* intervention the channel was not trained on (ablating the named concept matched to a relevant generative factor and asking whether the decision moves

by exactly the amount the true mechanism dictates), CANDOR is 81.7% directionally correct on relevant concepts and leaves the decision unchanged 99.6% of the time when an irrelevant concept is ablated. The necessity gap (the decision-flip rate when the atoms matched to the relevant concepts are ablated, versus an equal number of irrelevant atoms) is 77.3 percentage points (77.3% vs. 0.0%).

Objective ablations (Fig. 3). Each term removes a *distinct* guarantee, which is exactly why the conjunction, not any single part, is the contribution. Dropping the *leak-energy* term collapses concept recovery from 0.919 to 0.496: the code stops explaining the activation, even though δ stays small. Dropping the *causal* term keeps δ small but restores the leak’s causal power (leak-swap effect 0.0110, roughly double CANDOR’s 0.0054). This is behavioural mimicry without causal sufficiency, the exact failure a non-causal, reconstruction-only method cannot even detect. Dropping *anchoring* leaves within-run metrics intact but, across independently trained runs, concept-identity stability falls from 0.722 to 0.000, reproducing the SAE non-canonicity finding [Leask et al., 2025] and showing anchoring fixes it. The *faithfulness* term keeps the certificate tight (removing it raises δ modestly, to 0.0030). No single objective yields completeness, causal sufficiency, stable identity, and a tight certificate together; their conjunction does.

7.2 Dark-computation detection: a backdoor that cannot hide (Fig. 2)

We plant a rare trigger (4.5% of inputs), a direction *disjoint* from the named concepts, that flips the label. To use it, the model must route the computation through the leak, so the certificate should react. It does: δ on triggered inputs (mean 0.0040) is $\approx 6\times$ that on clean inputs (0.0006), and δ alone detects the backdoor with AUC 0.901 (the model itself reaches 100.0% accuracy, so the backdoor is fully operational). This is the central anti-Goodhart evidence: because δ is a measurement of behavioural divergence and not a learned monitor, computation that bypasses the named channel *cannot* be hidden while remaining functional: it either becomes a named concept (inspectable) or shows up as dark computation (a δ spike). This is the property that optimisation-fragile chain-of-thought monitors lack [Baker et al., 2025].

7.3 Composition with attention

The Legible Bottleneck is architecture-agnostic. To show it composes with self-attention, we train a one-layer Transformer with a Legible Bottleneck on its MLP hidden state on a task that *requires* attention: predict the maximum token of a length-8 sequence (10 classes, 10200 train / 1800 test), which can only be solved by aggregating information across positions. CANDOR’s deployed model reaches 99.9% and its leak-ablated legible replacement 99.9% (an interpretability tax of 0.00 percentage points), while certifying with mean $\delta = 0.0000$ and 100.0% model/explanation agreement. The Legible Bottleneck composes with self-attention without collapsing capability; a full study of larger algorithmic circuits (e.g. modular addition and induction) is left to future work.

7.4 A real pretrained LLM: certificate transfer and the limits of retrofitting

Does any of this survive contact with a real language model? We retrofit a Legible Bottleneck as a transcoder on one MLP layer (9 / 12) of the pretrained **GPT-2 small** (124M parameters), on real text, and measure the certificate by *actually splicing the channel into GPT-2* and running the real model (its own causal attention, the rest of its layers) on held-out text. We train two channels at matched sparsity ($m = 2048$, $k = 32$): a reconstruction-only **SAE** and a **CANDOR** channel (reconstruction + behavioural faithfulness + causal sufficiency).

There are two findings, and we report the second prominently because it is the more instructive. *First, the certificate machinery transfers to a real LLM:* splicing the channel and running GPT-2 yields a measurable, sound, per-input δ (total variation at the next token) of ≈ 0.1136 – 0.1299 at this sparsity, so the runtime certificate of §6 is not confined to toy models. *Second, in the frozen-retrofit regime the reconstruction-only SAE is the better channel:* it is modestly more behaviourally faithful ($\delta = 0.1136$ vs. CANDOR’s 0.1299), more causally inert (leak-swap 0.1319 vs. 0.1597), and reconstructs better (unexplained variance 0.375 vs. 0.486), and this ordering is stable across loss weightings, so it is structural, not a tuning artefact.

The reason, we hypothesised, is the thesis of this paper. On a *frozen* model the only lever that shrinks the leak is better reconstruction, precisely the SAE’s objective. CANDOR’s distinctive objectives (causal sufficiency, anchoring) assume the model can *reshape its own computation* to route through the legible channel; a frozen retrofit cannot, so they only spend capacity. The favourable tax/faithfulness frontier of Experiments 7.1–7.3 comes from training the model *with* the bottleneck, *by construction*, not from the dictionary alone. This delimits the contribution cleanly: CANDOR is not a drop-in replacement for a post-hoc SAE. It also yields a testable prediction: give the model room to reshape its computation and the behavioural and causal objectives should begin to pay. We test the smallest version of that prediction next.

Layer-wise fine-tuning: a dissociation, and the prediction does not yet hold. We unfreeze the spliced layer’s MLP (4.7M of 124M parameters) and fine-tune it jointly with the channel on the same text for 500 steps, under the full language-modelling loss (which anchors capability), in two conditions matched in every respect (unfrozen parameters, steps, batch, seed, learning rates; Appendix A) except the channel objective: reconstruction only (**FT-SAE**) versus reconstruction + behavioural faithfulness + causal sufficiency (**FT-CANDOR**). Two findings. *First, reconstruction and faithfulness dissociate.* Fine-tuning drives the control’s unexplained variance from 0.375 to 0.046 , an $\approx 8\times$ improvement, yet its certificate does not tighten ($\delta = 0.1282$, versus 0.1136 for the frozen retrofit; both fine-tuned models also adapt to the corpus, from 5.443 to ≈ 4.194 validation loss, so frozen-to-fine-tuned numbers compare channels on different underlying models, while the within-regime comparisons are exactly matched). Near-perfect reconstruction buys no behavioural faithfulness, which is precisely the gap between reconstruction-style objectives and guarantees that motivates this paper. *Second, the prediction does not hold at this scale.* At matched language-modelling loss (4.225 vs. 4.194), FT-CANDOR certifies *less* tightly than the control ($\delta = 0.1560$ vs. 0.1282) and its leak is less inert (0.1723 vs. 0.1485). Reshaping a single MLP under a joint objective, at this budget and weighting, is not the regime in which the behavioural and causal signals win; we report this negative result at the same standard as the positives (single seed, single weighting, every number measured and released). The by-construction gains of Experiments 7.1–7.3 have therefore *not yet* been reproduced on a pretrained model, by retrofitting or by local fine-tuning; the decisive experiment remains *full* by-construction training, co-training the network with its bottlenecks from the start, which exceeds the compute available for this study and which we name as the central open question (§9).

8 Related work

We state precedents prominently: *every component of CANDOR has antecedents*, and the novelty is their conjunction operating jointly and at runtime.

Superposition and the legibility substrate. [Elhage et al., 2022, 2021, Olah et al., 2020, Park et al., 2024, Hänni et al., 2024]. CANDOR inherits both the motivation for a sparse legible basis and the completeness limit it concedes (Rem. 1).




Figure 1: The interpretability-tax / legibility frontier. As the legible channel’s capacity k (active named concepts) grows past the task’s true concept count (dotted line), legible-model accuracy and concept recovery saturate while the certificate δ collapses: once the channel can carry the task’s concepts, legibility is nearly free.




Figure 2: Dark-computation detection. The per-input certificate $\delta(x)$ separates clean inputs from backdoor-triggered inputs whose computation bypasses the named channel (AUC 0.901). A bypass cannot be both functional and hidden.

Sparse dictionaries and their crisis. [Bricken et al., 2023, Cunningham et al., 2023, Templeton et al., 2024, Gao et al., 2024, Rajamanoharan et al., 2024, Bussmann et al., 2025, Lieberum et al., 2024, Chanin et al., 2024, Leask et al., 2025, Kantamneni et al., 2025]. CANDOR borrows the sparse typed channel and the reconstruction (completeness) objective; it departs by making the channel causal and certified rather than post-hoc, and by anchoring identity.

Transcoders, circuits, attribution. [Dunefsky et al., 2024, Marks et al., 2025, Ameisen et al., 2025, Lindsey et al., 2025, Wang et al., 2023, Conmy et al., 2023, Syed et al., 2023, Miller et al., 2024]. CANDOR borrows “the explanation is the forward pass” and error-node accounting; it departs by routing *into* the trained forward pass and bounding the error-node mass at runtime.

Causal faithfulness as an objective. [Geiger et al., 2021, 2022, 2024, Wu et al., 2023, Chan et al., 2022, Geiger et al., 2025, Makelov et al., 2024]. This is CANDOR’s core borrowed engine. CANDOR generalises IIT from aligning a *given* causal variable to co-training a typed, persistent, *complete* bottleneck plus a runtime certificate; it does not invent training-time causal faithfulness.

By-construction interpretability and its leakage. [Koh et al., 2020, Alvarez-Melis and Jaakkola, 2018, Espinosa Zarlenga et al., 2022, Chen et al., 2019, Böhle et al., 2022, van den Oord et al., 2017, Hewitt et al., 2023, Rudin, 2019, Mahinpei et al., 2021, Havasi et al., 2022, Locatello et al., 2019]. CANDOR is a concept bottleneck hardened with the guarantees this line proved necessary; *leakage is dark computation under another name*, and the certificate is what bounds it.

Certification, oversight, evaluation, governance. [Katz et al., 2017, Zhang et al., 2018, Xu et al., 2020, Przybysz et al., 2023, Ignatiev et al., 2019, Christiano et al., 2021, Roger et al., 2023, Lanham et al., 2023, Chen et al., 2025, Baker et al., 2025, Korbak et al., 2025, Jacovi and Goldberg, 2020, Karvonen et al., 2025, Amodei, 2025, Clymer et al., 2024, European Union, 2024]. CANDOR borrows the sound-over-approximation form and the safety-case framing, and supplies

Objective ablations: each term removes a distinct guarantee




Figure 3: Objective ablations: each term is load-bearing for a *distinct* guarantee. **Left:** the leak-energy term governs concept recovery (dropping it halves recovery). **Centre:** the causal term governs causal sufficiency (dropping it roughly doubles the leak’s causal effect, mimicry without sufficiency). **Right:** the anchoring term governs cross-run concept-identity stability (dropping it sends it to zero). No single objective delivers all three; the conjunction does.

the runtime certificate that monitoring and CoT approaches cannot. The runtime-certificate stance also continues a line of *certifying* runtime guardrails for deployed models, but moves the certificate from the model’s *outputs* to the model’s own *computation*.

9 Limitations and honest scope

- **Scale, and the retrofit/by-construction gap.** Our positive evidence is on synthetic ground-truth tasks and a one-layer Transformer, *not* a frontier LLM; we chose the regime where soundness can be *checked*. On the real pretrained GPT-2 (§7.4) the certificate *machinery* transfers, but neither a *frozen* retrofit nor single-layer joint fine-tuning favours CANDOR over a reconstruction-only SAE, and under fine-tuning an $\approx 8\times$ reconstruction improvement fails to tighten the certificate. The hypothesis that CANDOR’s gains require the model to reshape its computation through the channel has so far been tested only down to one layer’s MLP, one seed and one loss weighting, and it did not hold there; it remains open beyond that regime. The single most important next experiment is therefore *full* by-construction training of an LLM with its bottlenecks (plus a seed and weight sweep of the fine-tuning probe), followed by a frontier-scale evaluation (SAEBench, RAVEL, MIB, and a real backdoor-detection battery [Karvonen et al., 2025, Huang et al., 2024, Mueller et al., 2025]); these are future work, not present claims. Whether the favourable tax/legibility frontier persists at scale is the central open empirical question.
- **No completeness.** By Rem. 1 we do not and cannot claim the certificate reaches zero on arbitrary computation; we claim it is sound, measured, and trainable, and that a large δ is an honest abstention signal.
- **“Certificate” $\neq$ formal proof.** $\delta(x)$ is a sound *per-input measurement*, not a verified bound over all inputs. It is one-sided (cannot under-report divergence for the inputs evaluated) but says nothing, on its own, about inputs not run. A periodic exact certificate on a tractable channel fragment [Przybylski et al., 2023] and a worst-case bound-propagation audit [Xu et al., 2020] are the natural complements.
- **Anchoring tension.** Anchoring injects inductive bias to stabilise identity; pushed too hard it makes concepts the designer’s labels rather than the model’s. We report where CANDOR sits via the cross-run stability number rather than assuming the best case.

- **Concept placement.** On the planted task we place the concept layer where the generative factors are linearly present; identifying such sites automatically in a deep agent is unsolved, and computation that is irreducibly distributed across many sites will show up as leak, not as a clean concept.

10 Broader impact

This work is about transparency and oversight of capable AI. A runtime, sound, checkable bound on a model’s non-legible computation is exactly the kind of artefact a safety case for an autonomous agent needs [Clymer et al., 2024, Amodei, 2025] and that transparency regulation increasingly asks for [European Union, 2024]. The same mechanism is dual-use only in the weak sense that any interpretability tool is: it makes models easier to understand, audit, and correct. The honest risk we flag is *false assurance*: treating a low δ on a test distribution as a guarantee off-distribution. We therefore frame δ as a measurement with stated assumptions, demonstrate that it *rises* on bypassing computation rather than staying silent, and release everything needed to recompute it.

11 Conclusion

Interpretability has been practised as archaeology: build an opaque artefact, then dig. We have argued that the way toward interpretable AI systems is to stop digging, to make legibility an architectural invariant and emit the audit at inference time. The Legible Bottleneck is the single primitive that makes this possible: route a layer’s computation through a sparse, typed, anchored concept code, keep and *measure* the unexplained leak, and train the concepts to be causally sufficient. The resulting CANDOR architecture certifies each of its decisions with a sound, re-checkable bound on its own dark computation, recovers ground-truth concepts at near-zero tax, and surfaces a channel-bypassing backdoor that no post-hoc method bounds. It is not a solution to interpretable AI, and it does not claim completeness. It is a concrete, checkable mechanism for turning interpretability from a description one hopes is true into an invariant the model proves, one forward pass at a time.

Reproducibility statement

Every number in this paper is machine-generated: `experiments/run_all.py` (and `experiments/exp_gpt2.py` for §7.4) writes raw results to `results/*.json`, `scripts/paper_numbers.py` converts them into the macros used in the text (`paper/_numbers.tex`), and `scripts/make_figures.py` renders the figures. No number is hand-entered, and `make all` reproduces the paper end-to-end from a fresh checkout. All random seeds are fixed in the experiment scripts (Appendix A); the experiments run on commodity hardware (CPU is sufficient), and the GPT-2 experiment uses only the publicly available GPT-2 small weights and a public text corpus. Code, the reference implementation (`candorkit`), unit tests, and all measured results are available under an MIT license at <https://github.com/AnwarDebes/candor>.

References

- David Alvarez-Melis and Tommi S. Jaakkola. Towards robust interpretability with self-explaining neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.
- Emmanuel Ameisen, Jack Lindsey, Adam Pearce, Wes Gurnee, Nicholas L. Turner, Brian Chen, Craig Citro, Joshua Batson, et al. Circuit tracing: Revealing computational graphs in language models. *Transformer Circuits Thread*, 2025.

- Dario Amodei. The urgency of interpretability. <https://darioamodei.com/post/the-urgency-of-interpretability>, 2025.
- Bowen Baker, Joost Huizinga, et al. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv preprint arXiv:2503.11926*, 2025. OpenAI.
- Moritz Böhle, Mario Fritz, and Bernt Schiele. B-cos networks: Alignment is all we need for interpretability. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermy, Tom Conerly, Nicholas L. Turner, Cem Anil, Carson Denison, Amanda Askell, et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.
- Bart Bussmann, Patrick Leask, and Neel Nanda. BatchTopK sparse autoencoders. *arXiv preprint arXiv:2412.06410*, 2024.
- Bart Bussmann, Noa Nabeshima, Adam Karvonen, and Neel Nanda. Learning multi-level features with Matryoshka sparse autoencoders. In *International Conference on Machine Learning (ICML)*, 2025. arXiv:2503.17547.
- Lawrence Chan, Adrià Garriga-Alonso, Nicholas Goldowsky-Dill, Ryan Greenblatt, Jenny Nitishinskaya, Ansh Radhakrishnan, Buck Shlegeris, and Nate Thomas. Causal scrubbing: A method for rigorously testing interpretability hypotheses. Technical report, Redwood Research, Alignment Forum, 2022.
- David Chanin, James Wilken-Smith, Tomáš Dulka, Hardik Bhatnagar, and Adrià Garriga-Alonso. A is for absorption: Studying feature splitting and absorption in sparse autoencoders. *arXiv preprint arXiv:2409.14507*, 2024.
- David Chanin, Tomáš Dulka, and Adrià Garriga-Alonso. Feature hedging: Correlated features break narrow sparse autoencoders. *arXiv preprint arXiv:2505.11756*, 2025.
- Chaofan Chen, Oscar Li, Chaofan Tao, Alina Jade Barnett, Jonathan Su, and Cynthia Rudin. This looks like that: Deep learning for interpretable image recognition. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- Yanda Chen, Joe Benton, et al. Reasoning models don’t always say what they think. *arXiv preprint arXiv:2505.05410*, 2025. Anthropic Alignment Science.
- Paul Christiano, Ajeya Cotra, and Mark Xu. Eliciting latent knowledge: How to tell if your eyes deceive you. Technical report, Alignment Research Center, 2021.
- Joshua Clymer, Nick Gabrieli, David Krueger, and Thomas Larsen. Safety cases: How to justify the safety of advanced AI systems. *arXiv preprint arXiv:2403.10462*, 2024.
- Arthur Conmy, Augustine N. Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023.
- Jacob Dunefsky, Philippe Chlenski, and Neel Nanda. Transcoders find interpretable LLM feature circuits. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2406.11944.

- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Christopher Olah, et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021.
- Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah. Toy models of superposition. *Transformer Circuits Thread*, 2022. arXiv:2209.10652.
- Mateo Espinosa Zarlenga, Pietro Barbiero, Gabriele Ciravegna, Giuseppe Marra, Francesco Giannini, Michelangelo Diligenti, Zohreh Shams, Frederic Precioso, Stefano Melacci, Adrian Weller, Pietro Lio, and Mateja Jamnik. Concept embedding models: Beyond the accuracy-explainability trade-off. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- European Union. Regulation (EU) 2024/1689 (Artificial Intelligence Act), Article 13: Transparency and Provision of Information to Deployers. Official Journal of the European Union, 2024.
- Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2406.04093*, 2024.
- Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. Causal abstractions of neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- Atticus Geiger, Zhengxuan Wu, Hanson Lu, Josh Rozner, Elisa Kreiss, Thomas Icard, Noah D. Goodman, and Christopher Potts. Inducing causal structure for interpretable neural networks. In *International Conference on Machine Learning (ICML)*, 2022. arXiv:2112.00826.
- Atticus Geiger, Zhengxuan Wu, Christopher Potts, Thomas Icard, and Noah D. Goodman. Finding alignments between interpretable causal variables and distributed neural representations. In *Causal Learning and Reasoning (CLear)*, PMLR vol. 236, 2024.
- Atticus Geiger, Duligur Ibeling, Amir Zur, Maheep Chaudhary, Sonakshi Chauhan, Jing Huang, Aryaman Arora, Zhengxuan Wu, Noah D. Goodman, Christopher Potts, and Thomas Icard. Causal abstraction: A theoretical foundation for mechanistic interpretability. *Journal of Machine Learning Research*, 26(83), 2025. arXiv:2301.04709.
- Nicholas Goldowsky-Dill, Chris MacLeod, Lucas Sato, and Aryaman Arora. Localizing model behavior with path patching. *arXiv preprint arXiv:2304.05969*, 2023.
- Ole-Christoffer Granmo. The Tsetlin machine: A game theoretic bandit driven approach to optimal pattern recognition in propositional logic. *arXiv preprint arXiv:1804.01508*, 2018.
- Kaarel Hänni, Jake Mendel, Dmitry Vaintrob, and Lawrence Chan. Mathematical models of computation in superposition. In *ICML 2024 Mechanistic Interpretability Workshop*, 2024. arXiv:2408.05451.
- Marton Havasi, Sonali Parbhoo, and Finale Doshi-Velez. Addressing leakage in concept bottleneck models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- John Hewitt, John Thickstun, Christopher D. Manning, and Percy Liang. Backpack language models. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.

- Jing Huang, Zhengxuan Wu, Christopher Potts, Mor Geva, and Atticus Geiger. RAVEL: Evaluating interpretability methods on disentangling language model representations. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- Alexey Ignatiev, Nina Narodytska, and Joao Marques-Silva. Abduction-based explanations for machine learning models. In *AAAI Conference on Artificial Intelligence*, 2019.
- Alon Jacovi and Yoav Goldberg. Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness? In *Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 4198–4205, 2020.
- Subhash Kantamneni, Joshua Engels, Senthooan Rajamanoharan, Max Tegmark, and Neel Nanda. Are sparse autoencoders useful? a case study in sparse probing. In *International Conference on Machine Learning (ICML)*, 2025. arXiv:2502.16681.
- Adam Karvonen, Can Rager, Johnny Lin, Curt Tigges, Joseph Bloom, David Chanin, Yeu-Tong Lau, Eoin Farrell, Arthur Conmy, Callum McDougall, et al. SAEBench: A comprehensive benchmark for sparse autoencoders in language model interpretability. *arXiv preprint arXiv:2503.09532*, 2025.
- Guy Katz, Clark Barrett, David L. Dill, Kyle Julian, and Mykel J. Kochenderfer. Reluplex: An efficient SMT solver for verifying deep neural networks. In *Computer Aided Verification (CAV)*, 2017.
- Pang Wei Koh, Thao Nguyen, Yew Siang Tang, Stephen Mussmann, Emma Pierson, Been Kim, and Percy Liang. Concept bottleneck models. In *International Conference on Machine Learning (ICML)*, 2020.
- Tomek Korbak, Mikita Balesni, Elizabeth Barnes, Yoshua Bengio, Joe Benton, Marius Hobbhahn, et al. Chain of thought monitorability: A new and fragile opportunity for AI safety. *arXiv preprint arXiv:2507.11473*, 2025.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- Patrick Leask, Bart Bussmann, Michael Pearce, Joseph Bloom, Curt Tigges, Noura Al Moubayed, Lee Sharkey, and Neel Nanda. Sparse autoencoders do not find canonical units of analysis. *arXiv preprint arXiv:2502.04878*, 2025.
- Tom Lieberum, Senthooan Rajamanoharan, Arthur Conmy, Lewis Smith, Nicolas Sonnerat, Vikrant Varma, János Kramár, Anca Dragan, Rohin Shah, and Neel Nanda. Gemma Scope: Open sparse autoencoders everywhere all at once on Gemma 2. In *Proceedings of the 7th BlackboxNLP Workshop*, 2024. arXiv:2408.05147.
- Jimmy Lin, Chudi Zhong, Diane Hu, Cynthia Rudin, and Margo Seltzer. Generalized and scalable optimal sparse decision trees. In *International Conference on Machine Learning (ICML)*, 2020.
- Jack Lindsey, Wes Gurnee, Emmanuel Ameisen, Brian Chen, Adam Pearce, Nicholas L. Turner, Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, et al. On the biology of a large language model. *Transformer Circuits Thread*, 2025.
- Francesco Locatello, Stefan Bauer, Mario Lucic, Gunnar Rätsch, Sylvain Gelly, Bernhard Schölkopf, and Olivier Bachem. Challenging common assumptions in the unsupervised learning of disentangled representations. In *International Conference on Machine Learning (ICML)*, 2019.

- Anita Mahinpei, Justin Clark, Isaac Lage, Finale Doshi-Velez, and Weiwei Pan. Promises and pitfalls of black-box concept learning models. *arXiv preprint arXiv:2106.13314*, 2021.
- Aleksandar Makelov, Georg Lange, and Neel Nanda. Is this the subspace you are looking for? an interpretability illusion for subspace activation patching. In *International Conference on Learning Representations (ICLR)*, 2024.
- Samuel Marks, Can Rager, Eric J. Michaud, Yonatan Belinkov, David Bau, and Aaron Mueller. Sparse feature circuits: Discovering and editing interpretable causal graphs in language models. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2403.19647.
- Joseph Miller, Bilal Chughtai, and William Saunders. Transformer circuit faithfulness metrics are not robust. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2407.08734.
- Aaron Mueller, Atticus Geiger, Sarah Wiegrefe, et al. MIB: A mechanistic interpretability benchmark. In *International Conference on Machine Learning (ICML)*, 2025.
- Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter. Zoom in: An introduction to circuits. *Distill*, 2020. doi: 10.23915/distill.00024.001.
- Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, et al. In-context learning and induction heads. *Transformer Circuits Thread*, 2022. arXiv:2209.11895.
- Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry of large language models. In *International Conference on Machine Learning (ICML)*, 2024.
- Emilia Przybylska, Bimal Bhattarai, Cosimo Persia, Ana Ozaki, Ole-Christoffer Granmo, and Jivitesh Sharma. Verifying properties of Tsetlin machines. *arXiv preprint arXiv:2303.14464*, 2023.
- Senthoooran Rajamanoharan, Tom Lieberum, Nicolas Sonnerat, Arthur Conmy, Vikrant Varma, János Kramár, and Neel Nanda. Jumping ahead: Improving reconstruction fidelity with JumpReLU sparse autoencoders. *arXiv preprint arXiv:2407.14435*, 2024.
- Fabien Roger, Ryan Heiner, et al. Benchmarks for detecting measurement tampering. *arXiv preprint arXiv:2308.15605*, 2023.
- Cynthia Rudin. Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature Machine Intelligence*, 1(5):206–215, 2019.
- Aaquib Syed, Can Rager, and Arthur Conmy. Attribution patching outperforms automated circuit discovery. In *NeurIPS Workshop on Attributing Model Behavior at Scale (ATTRIB) / BlackboxNLP*, 2023. arXiv:2310.10348.
- Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Ameisen, Andy Jones, et al. Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024.
- Naftali Tishby and Noga Zaslavsky. Deep learning and the information bottleneck principle. In *IEEE Information Theory Workshop (ITW)*, 2015.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

- Aaron van den Oord, Oriol Vinyals, and Koray Kavukcuoglu. Neural discrete representation learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2211.00593.
- Shiqi Wang, Huan Zhang, Kaidi Xu, Xue Lin, Suman Jana, Cho-Jui Hsieh, and J. Zico Kolter. Beta-CROWN: Efficient bound propagation with per-neuron split constraints for complete and incomplete neural network verification. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- Zhengxuan Wu, Atticus Geiger, Thomas Icard, Christopher Potts, and Noah D. Goodman. Interpretability at scale: Identifying causal mechanisms in Alpaca. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Kaidi Xu, Zhouxing Shi, Huan Zhang, Yihan Wang, Kai-Wei Chang, Minlie Huang, Bhavya Kailkhura, Xue Lin, and Cho-Jui Hsieh. Automatic perturbation analysis for scalable certified robustness and beyond. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. In *International Conference on Learning Representations (ICLR)*, 2024.
- Huan Zhang, Tsui-Wei Weng, Pin-Yu Chen, Cho-Jui Hsieh, and Luca Daniel. Efficient neural network robustness certification with general activation functions. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

A Experimental details

This appendix records every architecture, hyperparameter, and seed, as fixed in the released code. Optimisation is plain Adam throughout; no learning-rate schedules, weight decay, data augmentation, or early stopping are used anywhere. The anchoring reference $\bar{W}$ is an exponential moving average of the decoder with momentum 0.99, updated once per step.

A.1 Planted concepts (§7.1)

Data. 20,000 samples; $K = 12$ independent Bernoulli concepts with density 0.3, mixed linearly into 48 dimensions with additive Gaussian noise ($\sigma = 0.1$); the label counts the active concepts among a fixed relevant subset of size 5 (6 classes); generation seed 0; 80/20 train/validation split (seed 0).

Model. CANDOR-MLP: the concept layer sits on a frozen, identity-initialised embedding of the input; dictionary $m = 48$, $k = 8$; task head with two hidden layers of width 128.

Training. 2200 steps, Adam, learning rate 2×10^{-3} , batch size 512, seed 0. Loss weights: $\alpha = 0.5$ (legible task), $\beta = 1.0$ (faithfulness), $\lambda = 1.0$ (leak energy), $\mu = 0.5$ (causal sufficiency), $\gamma = 10^{-2}$ (anchoring), plus an L_1 code penalty of 2×10^{-3} .

Baseline. An opaque MLP of matched shape (width 128, two hidden layers) trained for the same 2200 steps (learning rate 2×10^{-3} , batch 512), followed by a post-hoc TopK SAE ($m = 48$,

$k = 8$) fitted to its embedding for 2,000 steps (Adam, learning rate 10^{-3} , batch 1,024, MSE reconstruction, decoder columns renormalised each step).

Causal tests. Atoms are matched to generative concepts by Hungarian assignment on absolute Pearson correlation. The *necessity* test ablates (zeroes) the atoms matched to the relevant concepts and compares the legible decision-flip rate against ablating an equal number of irrelevant atoms. The *directional* test ablates one relevant concept’s atom at a time, on the inputs where that concept is active, and checks that the legible decision moves by exactly the amount the ground-truth mechanism dictates (one class down, for the sum task), and symmetrically that ablating an irrelevant concept’s atom leaves the decision unchanged. Both tests run on up to 1,500 held-out validation inputs; neither intervention appears in the training objective.

Ablations. Each ablation retrains from the same initialisation (seed 0) with one weight zeroed: no-faithfulness ($\beta = 0$), no-leak ($\lambda = 0$), no-causal ($\mu = 0$), no-anchoring ($\gamma = 0$); remaining weights $\beta = 1.0$, $\lambda = 1.0$, $\mu = 0.5$, $\gamma = 10^{-2}$ as applicable, L_1 penalty 10^{-3} .

Frontier sweep (Fig. 1). The tax/legibility frontier retrains the same architecture at channel capacities $k \in \{1, 2, 3, 4, 6, 8, 12\}$ for 1,600 steps each (learning rate 2×10^{-3} , batch 512, seed 0) with the default loss weights ($\alpha = 0.5$, $\beta = 1.0$, $\lambda = 0.3$, $\mu = 0.3$, $\gamma = 10^{-2}$, L_1 10^{-3}), on the same data and split as the main run.

Stability protocol. Three training runs (seeds 11, 22, 33) per condition; the *anchored* condition uses anchoring weight 10^{-2} and a shared dictionary initialisation (the anchor’s reference point), the *free* condition uses no anchoring and per-seed initialisation; remaining weights at their defaults ($\alpha = 0.5$, $\beta = 1.0$, $\lambda = 0.3$, $\mu = 0.3$, L_1 10^{-3}). Reported stability is the mean pairwise concept-identity agreement across runs.

Backdoor (§7.2). 20,000 samples with $K = 12$ concepts of which 3 are relevant, density 0.3, noise 0.1; the base task is the parity of the relevant concepts (binary). The trigger is a fixed unit direction in 4 extra input dimensions (total input dimension 52), disjoint from the concept mixing directions by construction, added with magnitude 3.0 against background noise $\sigma = 0.1$; it is planted in a nominal 5% of inputs (measured 4.5% of the validation split) and flips the label. The model (same architecture) is trained with the default loss weights for 2200 steps; detection AUC uses the per-input $\delta(x)$ as the score, with no access to trigger labels.

A.2 Sequence task (§7.3)

Data. 12,000 sequences of length 8 over a vocabulary of 10 tokens (generation seed 0); the label is the maximum token (10 classes); 85/15 split (10200 train / 1800 test).

Model. One-layer Transformer: $d_{\text{model}} = 96$, 4 attention heads, MLP width 256, with a Legible Bottleneck ($m = 128$, $k = 8$) on the MLP hidden state.

Training. 1,200 steps, Adam, learning rate 10^{-3} , batch size 256, seed 0; loss weights $\alpha = 0.5$, $\beta = 1.0$, $\lambda = 0.5$, $\mu = 0.3$, $\gamma = 10^{-2}$, L_1 penalty 10^{-3} ; trained on CPU.

A.3 GPT-2 retrofit (§7.4)

Setup. GPT-2 small (124M parameters, frozen throughout); a transcoder-style Legible Bottleneck reads the MLP input and predicts the MLP output at layer 9 of 12, spliced into the live model via a forward hook. Corpus: the public TinyShakespeare text, tokenised with the GPT-2 tokeniser into 1,400 sequences of length 24; 85/15 split (1,190 train / 210 validation sequences).

Channels. Both channels use $m = 2048$, $k = 32$, unit-norm decoder columns, and identical training budgets: 250 steps, Adam, learning rate 2×10^{-3} , batch size 12, seed 0, on CPU. The *SAE* channel minimises normalised MSE reconstruction of the MLP output only. The *CANDOR* channel adds the behavioural faithfulness term (KL between the spliced legible model’s next-token distribution and the cached full-model distribution) and the causal-sufficiency term (total variation against the leak-swapped splice), each with weight 1.0.

Metrics. δ is the mean total-variation distance between the full and legible-spliced next-token distributions on held-out text; the leak-swap effect is the same quantity against the leak-swapped splice; unexplained variance is residual energy over output energy; the loss-recovered fraction compares the legible splice’s language modelling loss against the full model and a zero-ablated MLP.

A.4 GPT-2 layer-wise fine-tuning (§7.4)

Setup. Same model, layer, corpus, tokenisation, and 85/15 split as the frozen retrofit. The MLP of the spliced layer (4.7M parameters) is unfrozen; all other model parameters stay frozen. The model is kept in evaluation mode throughout, so dropout is off (evaluation mode does not block gradients).

Conditions. Both conditions train the channel ($m = 2048$, $k = 32$) jointly with the unfrozen MLP for 500 steps (Adam; channel learning rate 2×10^{-3} , MLP learning rate 10^{-4} ; batch 12; seed 0) under the full language-modelling cross-entropy over all positions, which anchors capability, plus the normalised reconstruction loss with a *live* (non-detached) target, so in both conditions the MLP may reshape its output toward reconstructability. The FT-CANDOR condition adds the behavioural faithfulness term (KL between the spliced legible model’s next-token distribution and the current full model’s) and the causal-sufficiency term (total variation against the leak-swapped splice), each with weight 1.0; FT-SAE adds nothing. Matched step counts mean FT-CANDOR spends roughly three forward passes per step to FT-SAE’s one, as in the frozen comparison. Evaluation is identical to the frozen experiment, and the original model’s validation loss (5.443) is reported as the capability reference. This is a single-seed, single-weighting probe, stated as such in §7.4.

A.5 Compute

All experiments run on a single commodity machine; no data-centre resources are used. Each `results/*.json` file records the device and wall-clock time of the run that produced it.